



SAPIENZA
UNIVERSITÀ DI ROMA

Demonstrating real advantage of machine-learning-enhanced Monte Carlo for combinatorial optimization

Francesco Zamponi

Ispra, September 10, 2026



Luca Maria Del Bono, Federico Ricci-Tersenghi
PNAS (2026)

Federico Celauro
(in preparation)

Motivations

Goal: generate equilibrium samples from a Boltzmann distribution $P_B(\sigma) = \frac{e^{-\beta H(\sigma)}}{Z}$

For $\beta \rightarrow \infty$ find the minimum of $H(\sigma)$ - combinatorial optimization


Universal strategy: local MCMC + parallel tempering + population annealing +

But: there are hard-to-sample problems, among which glasses, NP-hard optimization...

Theorem: slow local decorrelation is due to spatial correlations between variables

Montanari, Semerjian (2006)

Propose 'smart moves' to update clusters of correlated variables?



A universal solution?

Global MCMC scheme - 'learn' a model that generates the right moves

1. Propose a new configuration σ_{new} sampled from a generative model $P_A(\sigma)$
2. Accept the move with Metropolis-Hastings rate

$$\text{Acc} [\sigma_{old} \rightarrow \sigma_{new}] = \min \left[1, \frac{P_B(\sigma_{new}) \times P_A(\sigma_{old})}{P_B(\sigma_{old}) \times P_A(\sigma_{new})} \right]$$

NB: similar to importance sampling - but can (and must!) be combined with local MCMC

Gabri , Rotskoff, Vanden-Ejiden (2021)
Del Bono, Ricci-Tersenghi, FZ (PRE 2025)

It will work if:

1. The reweighing $W(\sigma) = P_B(\sigma)/P_A(\sigma)$ does not fluctuate wildly
 2. States from $P_A(\sigma)$ cover well the support of $P_B(\sigma)$
- } $\Leftrightarrow D_{KL}(P_B || P_A)$ is small

$$D_{KL}(P_B || P_A) = \langle \log W(\sigma) \rangle_{P_B} \ll N$$

Learn P_A by minimizing $D_{KL}(P_B || P_A)$

The dangers of improper benchmarking

Article | Published: 21 April 2022

Combinatorial optimization with physics-inspired graph neural networks

Martin J. A. Schuetz , J. Kyle Brubaker  & Helmut G. Katzgraber 

Nature Machine Intelligence 4, 367–377 (2022) | [Cite this article](#)

12k Accesses | 179 Citations | 69 Altmetric | [Metrics](#)

 [Matters Arising](#) to this article was published on 30 December 2022

 [Matters Arising](#) to this article was published on 30 December 2022

Matters Arising | Published: 30 December 2022

Inability of a graph neural network heuristic to outperform greedy algorithms in solving combinatorial optimization problems

Stefan Boettcher 



Nature Machine Intelligence 5, 24–25 (2023) | [Cite this article](#)

Matters Arising | Published: 30 December 2022


Modern graph neural networks do worse than classical greedy algorithms in solving combinatorial optimization problems like maximum independent set

Maria Chiara Angelini  & Federico Ricci-Tersenghi

Nature Machine Intelligence 5, 29–31 (2023) | [Cite this article](#)




Algorithms



- Identical implementation on a single GPU (using MADE for GA)
- Validated benchmark instances of Quadratic Unconstrained Binary Optimization: random instances of 3dEA, check instance-to-instance fluctuations
- Fair wall-clock time comparison
- Check scalability with system size

Two instances of 3dEA




One 'easy', one 'hard' instance with $L = 10, N = 1000$, success over 50 runs

GA₀ does not work - need for local moves

Easy instance: PA outperforms GA

Hard instance: PA ~ GA, but GA is much sharper


Many instances of 3dEA



200 instances with $L = 10, N = 1000$, success over 10 runs

GA is more robust to instance-to-instance fluctuations
 GA is more efficient for very hard instances


Larger instances of 3dEA



10 instances with $L = 14, N = 2744$, success over 10 runs

GA is more efficient and more robust for large instances

Overlap probability distribution



A successful run on the 'hard' instance with $L = 10, N = 1000$

SA relies on rare events - one chain finds the ground state

PA relies on improved thermalization by reweighting - needs a large population

GA does the same more efficiently via the generative model

Notebook and datasets for benchmarking

Notebook



tinyurl.com/2ehkjuy8

Spin glass benchmarks



ocadni.github.io/spin-glass-benchmarks

Physics Reports 523 (2013) 127–205



Contents lists available at [SciVerse ScienceDirect](#)

Physics Reports

journal homepage: www.elsevier.com/locate/physrep



The quantum adiabatic algorithm applied to random optimization problems: The quantum spin glass perspective

V. Bapst^a, L. Foini^b, F. Krzakala^c, G. Semerjian^a, F. Zamponi^{a,*}

^a LPT, École Normale Supérieure, UMR 8549 CNRS, 24 Rue Lhomond, 75005, France

^b LPTHE, UPMC Paris 06, 4 Place Jussieu, 75252 Paris Cedex 05, France

^c ESPCI ParisTech, CNRS UMR 7083 Gulliver, 10 rue Vauquelin, 75005 Paris, France

Some progress



- Architectures scale well with $O(N)$ parameters. 
- Global annealing is robust against instance-to-instance fluctuations 
- It can be used to find the ground state efficiently in 2dEA and 3dEA 
- Not transferable: one has to learn a specific model for each sample of the problem. 
- What about really hard problems with exponentially many states? 

Ciarella, Trinquier, Weigt, FZ (2023)



- Better understanding of the role of correlations during annealing
- Use entirely different architectures, e.g. diffusion models
- Use the couplings as input to the GNN. Transferable architectures?
- Study harder problems, e.g. lattice glass or XORSAT
- Construct good benchmarks for further testing

Thank you for your attention!



Simone Ciarella, Jeanne Trinquier, Martin Weigt
Machine Learning: Science and Technology (2023)

Luca Maria Del Bono, Federico Ricci-Tersenghi
Machine Learning: Science and Technology (2025)
Physical Review E (2025)
PNAS (2026)